

Figure 2(a): End-to-End Transfer Learning Computational & Tensor Pipeline

Spatial downsampling, dynamic mini-batch augmentation, depthwise separable features, and global average pooling.




Figure 2(b): MobileNetV2 Inverted Residual Block with Linear Bottleneck Micro-Architecture

Pointwise expansion (1×1 Conv2D) -> Depthwise spatial filtering (3×3 Dwise) -> Linear projection bottleneck (1×1 Conv2D).




Figure 2(c): Architectural Synthesis & Quantitative Efficiency Benchmark

Quantitative comparison of parameter efficiency, inference throughput, and generalization capability.

Model Configuration	Total Parameters	Trainable Weights	Test Acc	Latency record
Custom Scratch CNN (Scratch Baseline)	6,944,449	6,943,489 (99.99%)	84.15%	12.1 ms aggregate
MobileNetV2 Phase 1 (Frozen Backbone)	2,586,177	328,193 (12.7%)	96.31%	18.4 ms aggregate
MobileNetV2 Ablation (No Augmentation)	2,586,177	1,842,433 (71.0%)	95.12%	18.4 ms aggregate
MobileNetV2 Fine-Tuned (Top 20 Layers FT)	2,586,177	986,369 (38.0%)	96.80%	18.4 ms aggregate
MobileNetV2 Final (Top 40 FT + Augmentation)	2,586,177	1,842,433 (71.0%)	97.47%	18.4 ms aggregate

Key Architectural Highlights

62.8% Parameter Reduction
GAP eliminates massive 2.65M dense parameter bottleneck.

Sample Efficiency
ImageNet priors enable convergence with only 18,000 training images.

Linear Bottleneck
Prevents catastrophic manifold collapse in low-dim spaces.

Real-time Inference
18.4 ms aggregate record; hardware metadata unavailable.