



Panel (a): End-to-End Transfer Learning & Tensor Dataflow Pipeline (MobileNetV2 Backbone)

Sequential spatial downsampling, dynamic online augmentation, depthwise feature factorization, and GAP classification head



Panel (b): MobileNetV2 Inverted Residual Block with Linear Bottleneck (Micro-Architecture)

Expansion (1x1 Conv2D) → Depthwise Conv (3x3 Dwise) → Linear Projection (1x1 Conv2D) with Identity Residual Shortcut



Panel (c): Architectural Efficiency, Computational Complexity & Benchmark Synthesis

Quantitative performance, parameter efficiency, and latency comparison between Custom Scratch CNN and MobileNetV2 Transfer Learning

Architectural Dimension	Custom Scratch CNN	MobileNetV2 (Transfer Learning)
Base Backbone Layers	4 Conv Blocks (Scratch)	154 Inverted Residuals (ImageNet)
Total Model Parameters	4,817,537 (~4.82 M)	2,586,177 (~2.59 M) [-46.3%]
Feature Map Transition	Flatten: 20x20x256 → 102,400	GAP: 5x5x1280 → 1,280 [-98.7%]
Dense Head Parameters	26,214,657 (Extreme Overhead)	328,193 (Compact & Regularized)
Computational Primitives	Standard 2D Convolution	Depthwise Separable (8.5x FLOPs Reduction)
Validation Accuracy	71.49% (High Overfitting Gap)	97.47% (Phase 2 Fine-Tuned)
Commercial Precision	71.85% (Unacceptable for Launch)	97.59% (Exceeds 90-95% Target)
Edge Latency (INT8)	64.2 ms / frame	15.3 ms / frame (65.4 FPS)

Key Architectural Highlights

8.5x FLOPs Reduction
Depthwise Separable Convolution

-98.7% Head Parameters
Global Average Pooling vs. Flatten

97.59% Dog Precision
Two-Phase Fine-Tuned ImageNet Prior

65.4 FPS INT8 Latency
Real-Time Edge Embedded Streaming