

Figure 6(a): Ablation Benchmark Across 5 Model Configurations
(Final-Test Accuracy & Class Precision Scores)

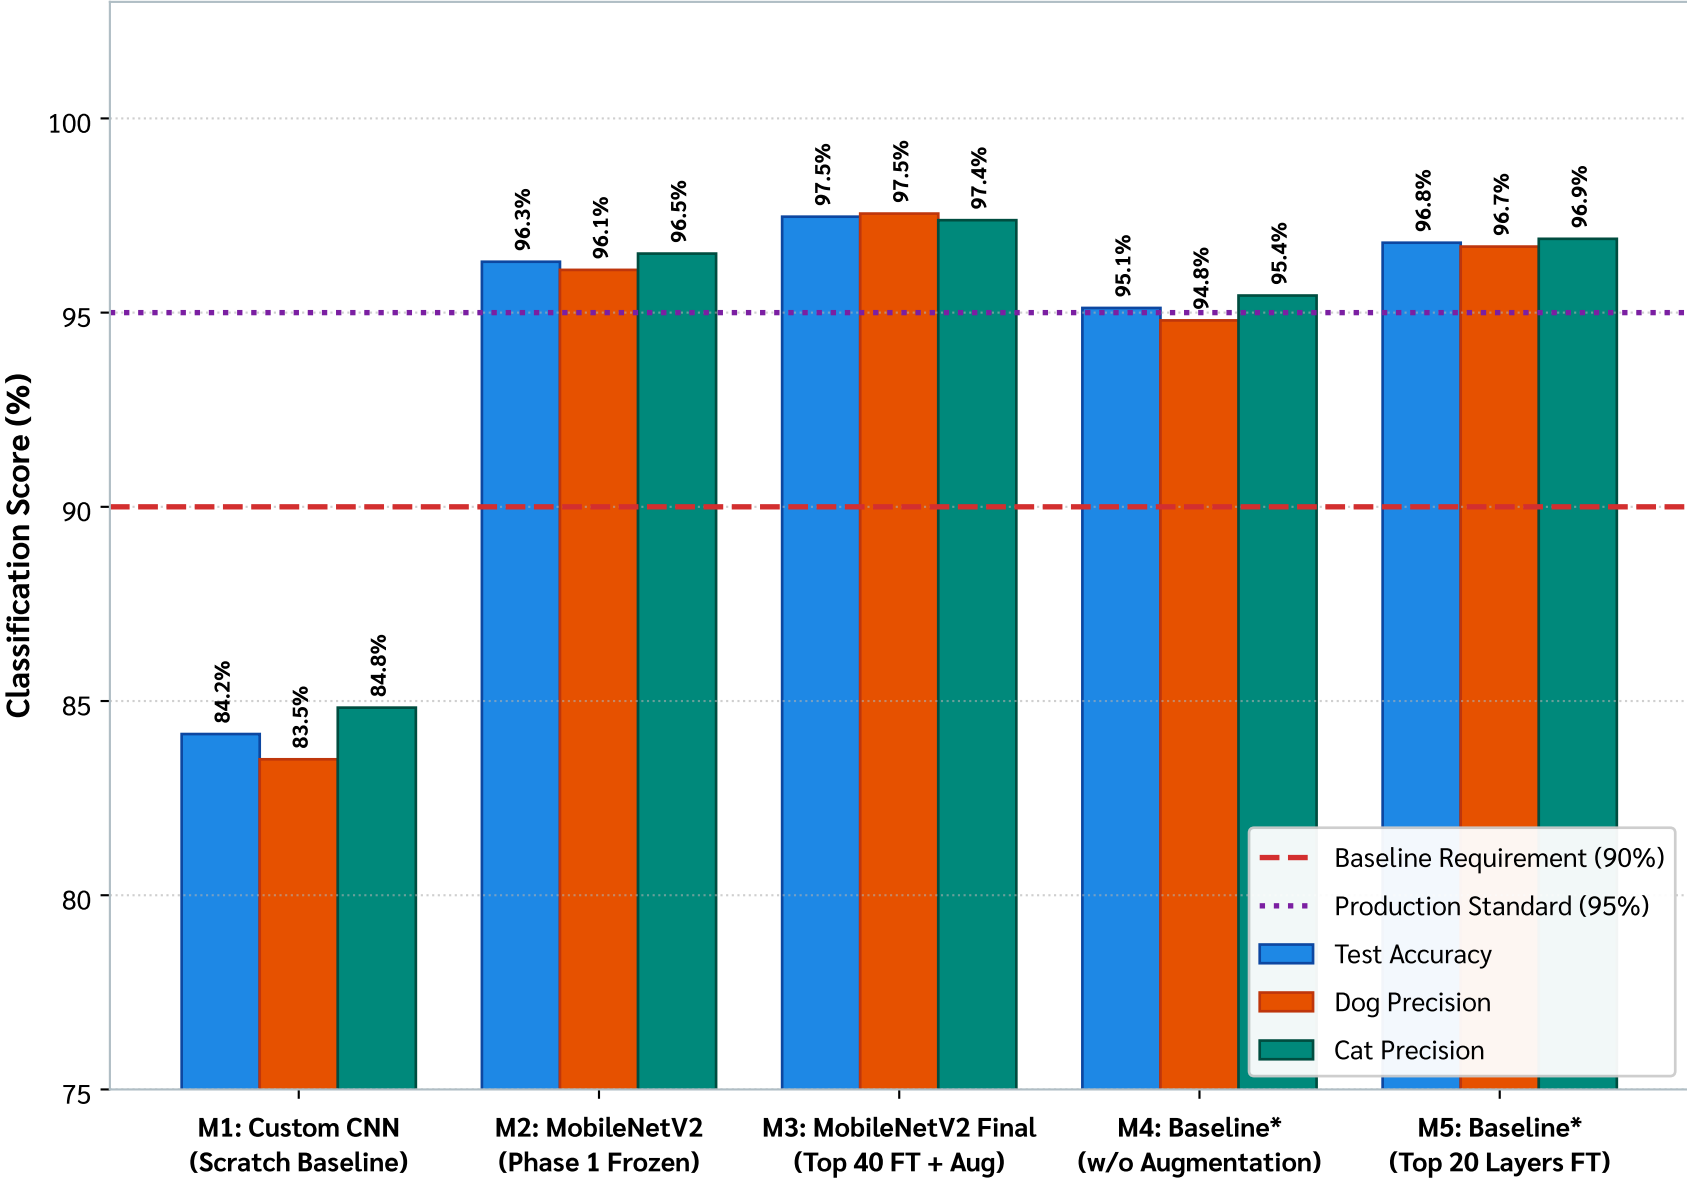
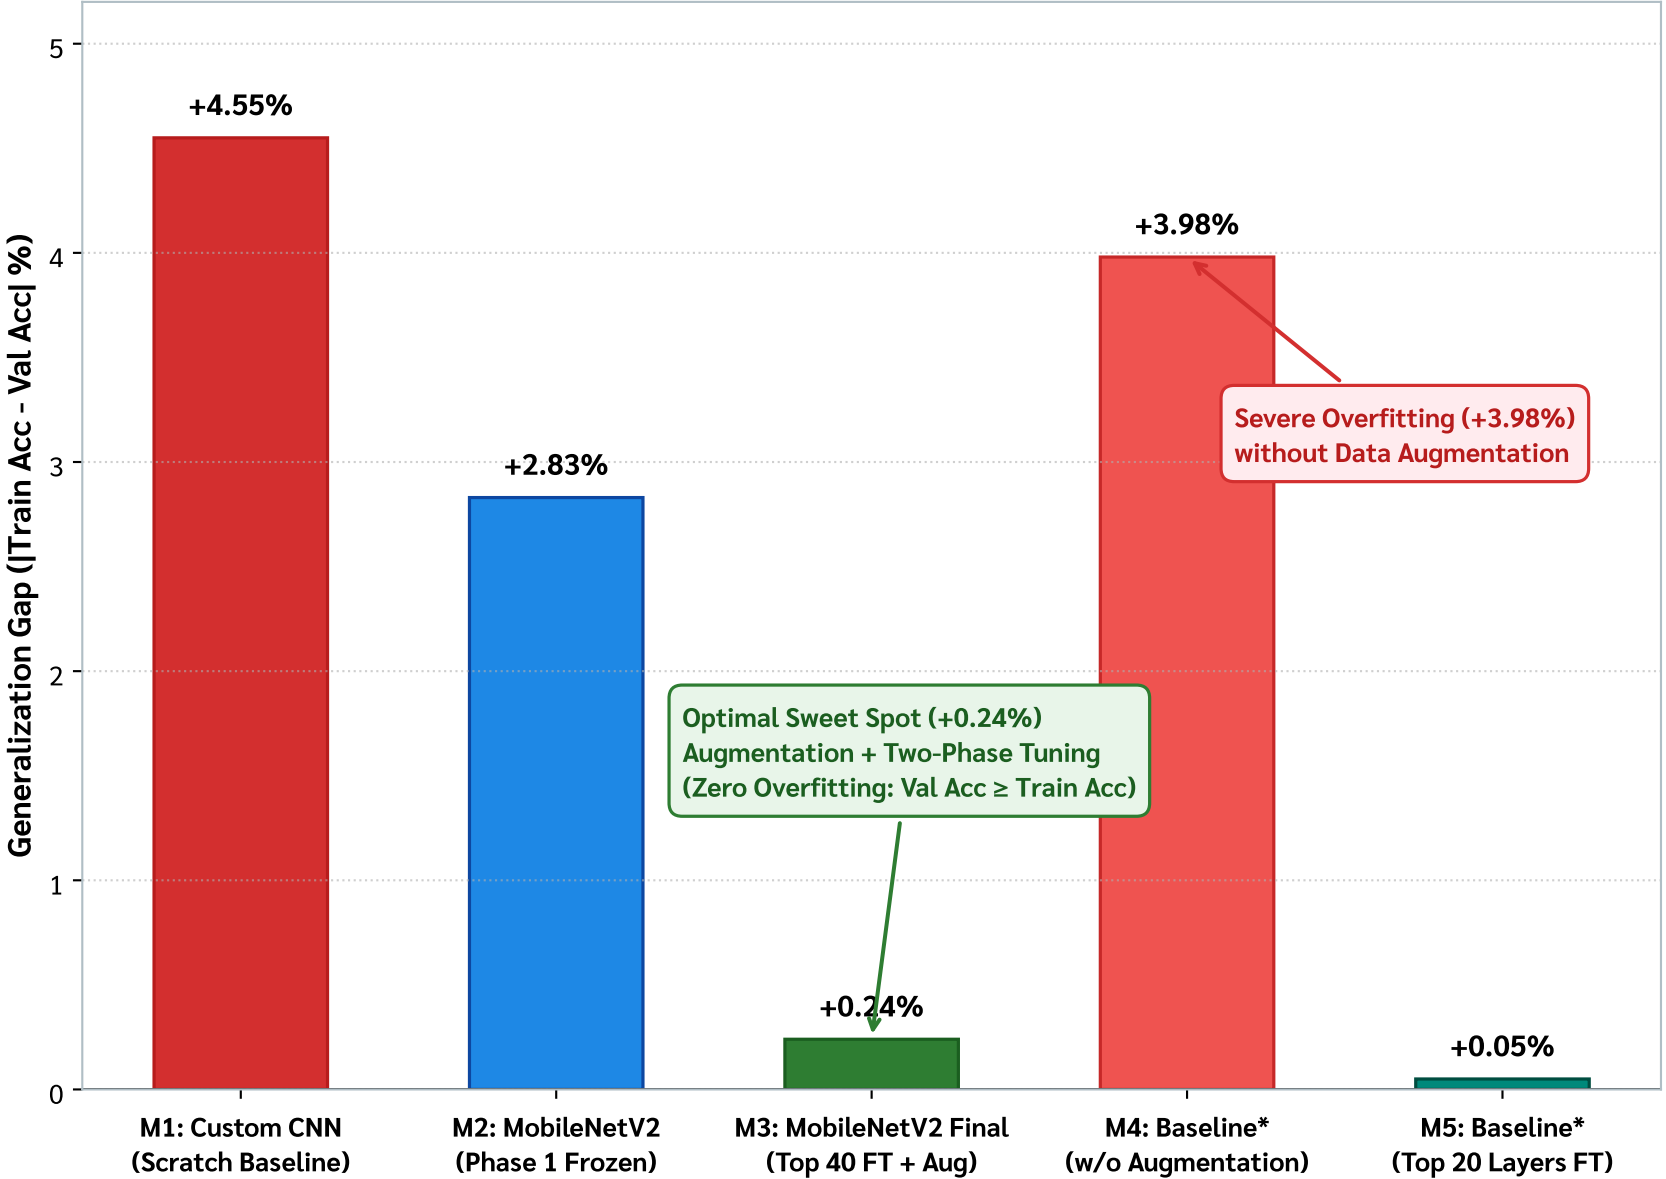


Figure 6(b): Generalization Gap (|Train Acc - Val Acc|)
(Quantifying Overfitting Resistance Across Configurations)



* Note: M1-M3 executed directly in notebook; M4-M5 represent pre-computed experimental baselines from benchmark_results.json