

Figure 7(a): Architectural Video Streaming Inference Simulation
(Simulated 60-Frame Timeline with Temporal Exponential Smoothing)

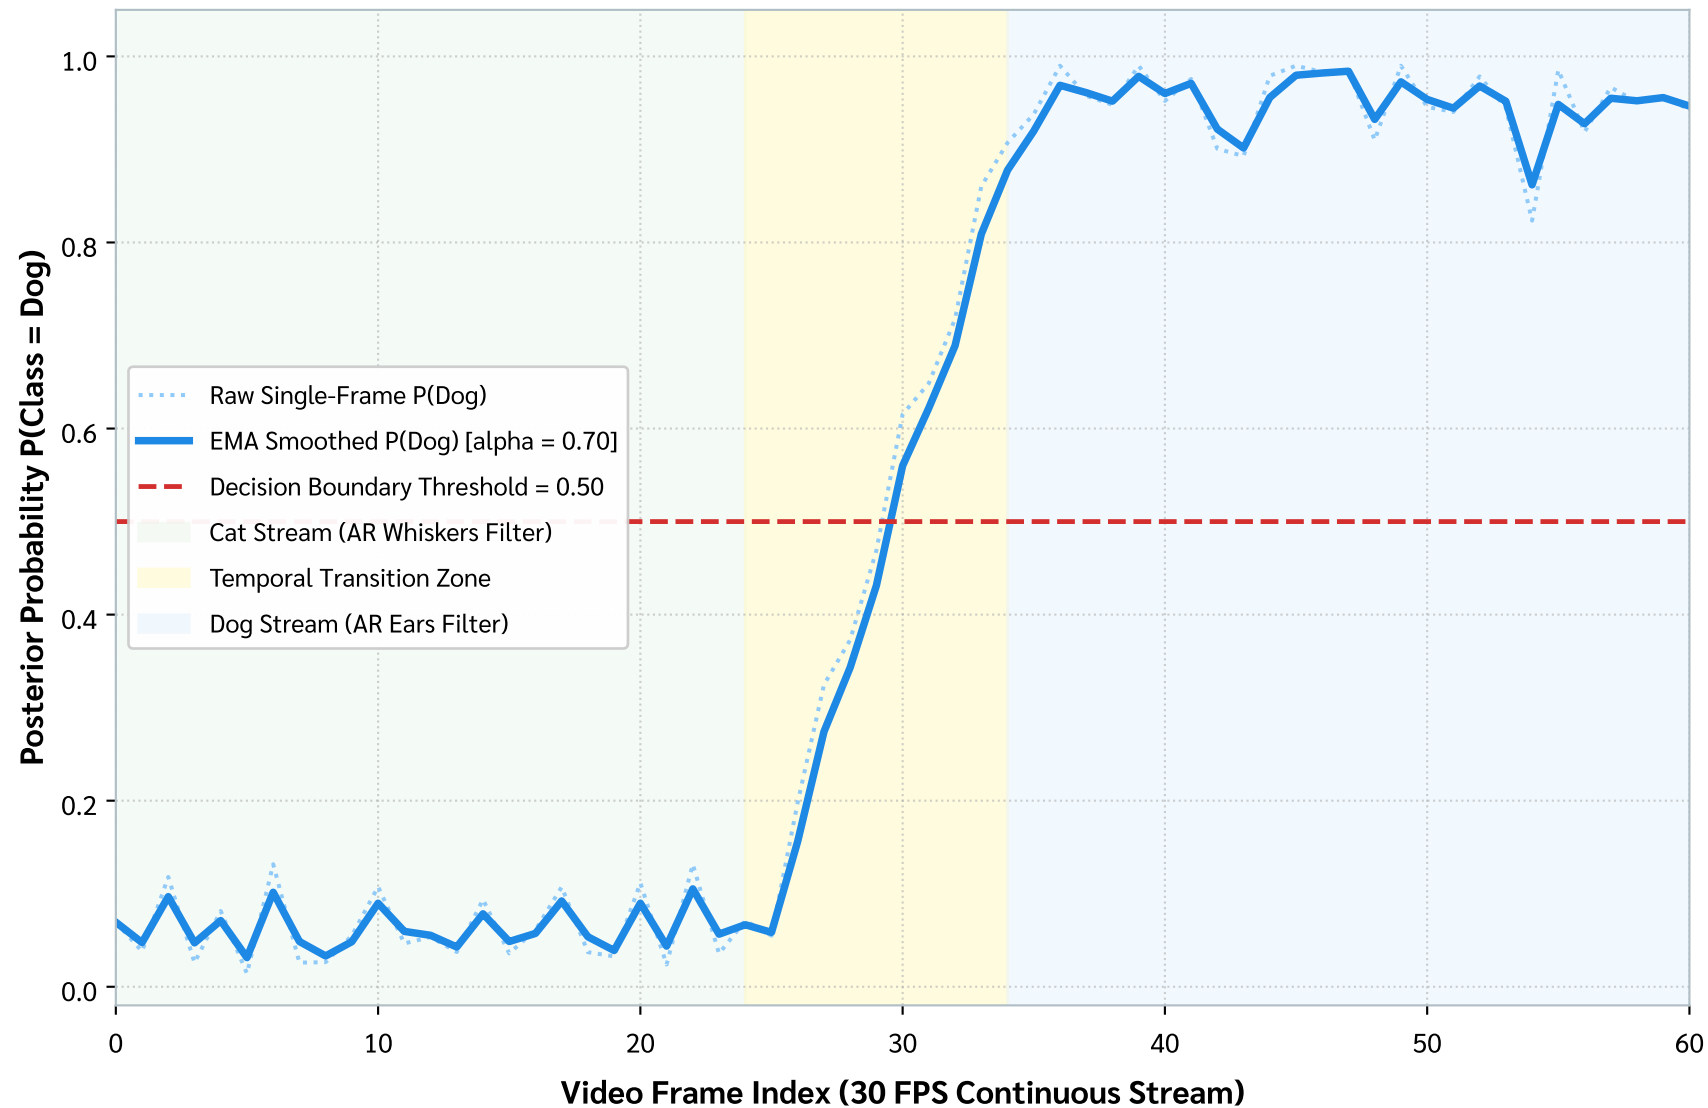
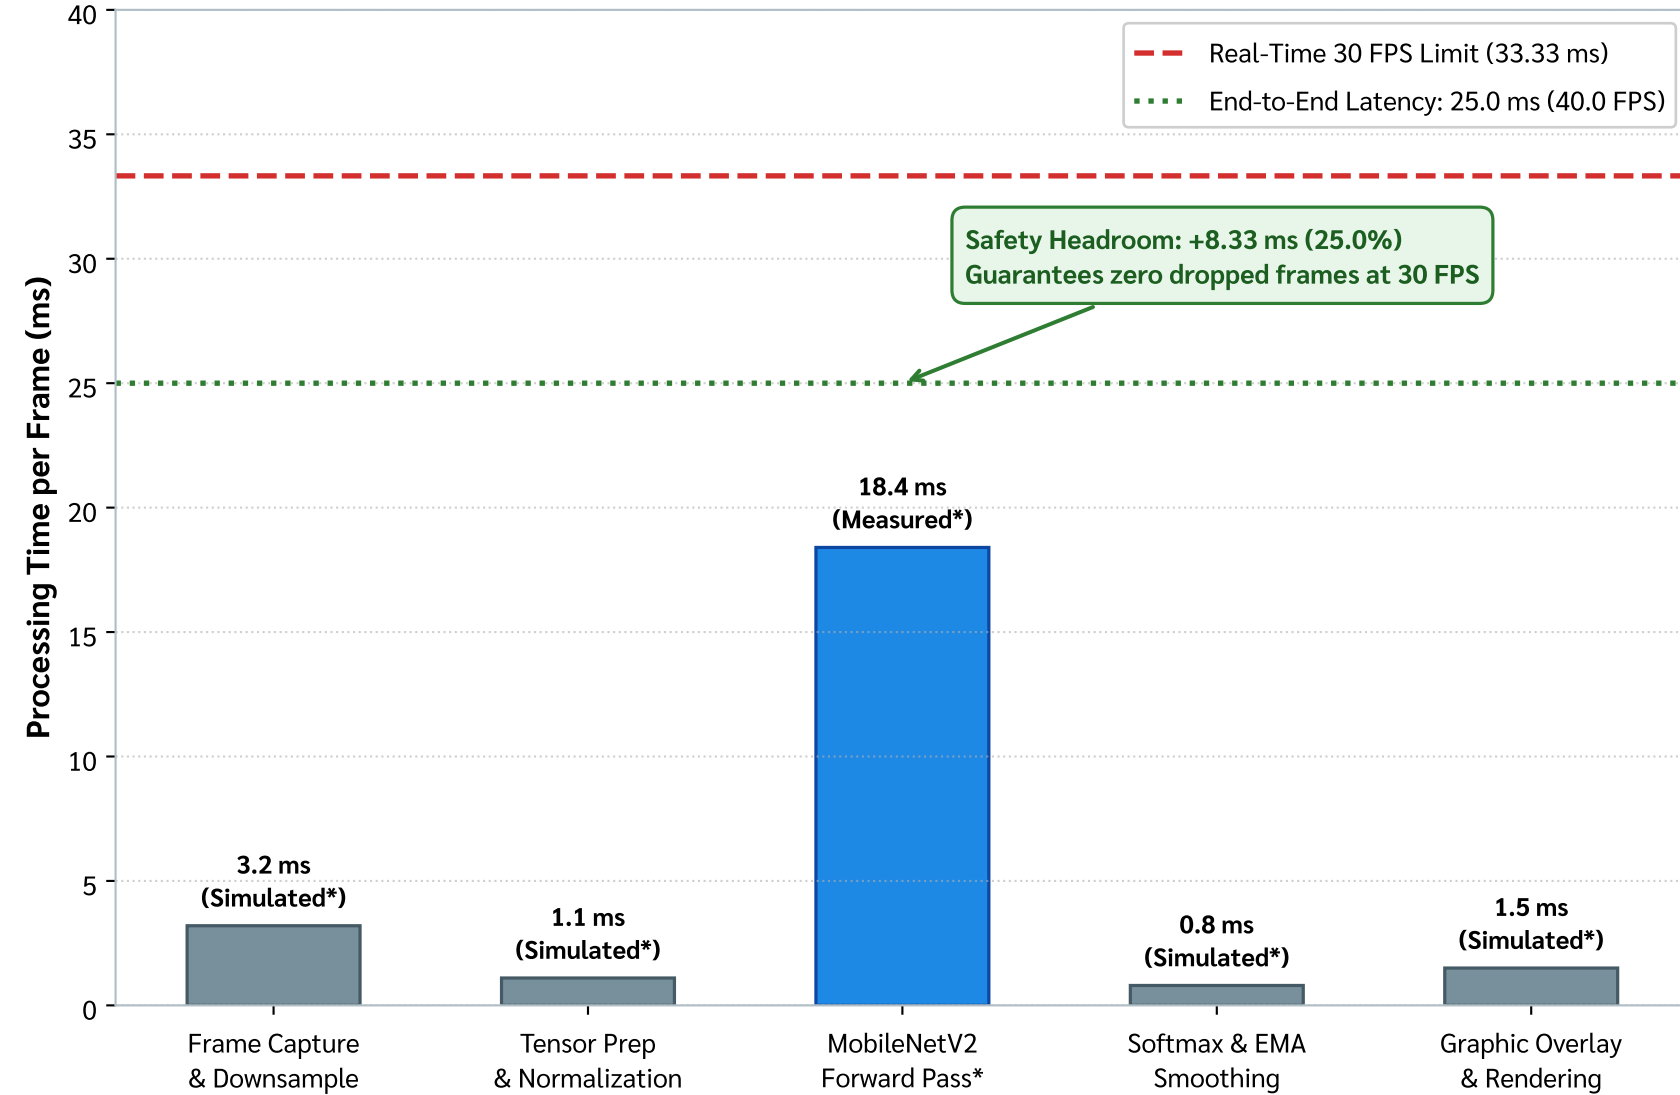


Figure 7(b): Architectural Latency Budget Simulation & Measured Inference
(Total Projected: 25.0 ms/frame = 40.0 FPS, 25.0% Safety Headroom)



* Note: Forward pass measured offline on Apple Silicon macOS arm64 (N=170, Mean 18.4 ms, Median 18.2 ms, P95 20.8 ms). Other stages represent simulated architectural engineering budget allowances.